

title: CVE-2016-4669分析与调试

date: 2016-11-07 14:57:00

categories: [CVE, OS X]

tags: [XNU,CVE,POC]

---

## 0x00 摘要

---

本文记录了对[CVE-2016-4669](#)的POC的调试中遇到的问题，以及相关知识的整理，该漏洞的原始报告在[这里](#)。原始报告中的内容，本文不在复述。

## 0x01 基础知识

---

### 1.1 MIG

---

MIG 是 Mach 系统中使用的一种自动生成代码的脚本语言，以 `.def` 结尾。通过工具生成的代码分为 `xxx.h`, `xxxClient.c` 和 `xxxServer.c` 三个部分，在编译应用层程序时，和 `xxxClient.c` 文件一起编译，使用自动生成的代码。这里就是 `poc` 中的 `taskUser.c` 和 `task.h`。

相关的细节可以查看《Mac OS X Internals: A Systems Approach》一书的**Section 9.6**中的描述。

### 1.2 内核中的内存管理

---

描述内核中堆内存的管理相关内容可以参考这个slides，比较简单明了的说清楚了内核中内存的基本结构[iOS 10 Kernel Heap Revisited](#)。

## 0x02 调试过程

---

### 2.1 core文件分析

---

在运行 `POC` 之后系统崩溃，查看崩溃的调用栈。

```

1 (lldb) bt
2 * thread #1: tid = 0x0000, 0xffffffff80049c0f01 kernel`hw_lock_to + 17, stop
  reason = signal SIGSTOP
3   * frame #0: 0xffffffff80049c0f01 kernel`hw_lock_to + 17
4     frame #1: 0xffffffff80049c5cb3 kernel`usimple_lock(l=0xdeadbeefdeadbef7) +
  35 at locks_i386.c:365 [opt]
5     frame #2: 0xffffffff80048c991c kernel`ipc_port_release_send [inlined]
  lck_spin_lock(lck=0xdeadbeefdeadbef7) + 44 at locks_i386.c:269 [opt]
6     frame #3: 0xffffffff80048c9914
  kernel`ipc_port_release_send(port=0xdeadbeefdeadbeef) + 36 at ipc_port.c:1567
  [opt]
7     frame #4: 0xffffffff80048e22d3 kernel`mach_ports_register(task=
  <unavailable>, memory=0xffffffff800aad4270, portsCnt=3) + 547 at ipc_tt.c:1097
  [opt]
8     frame #5: 0xffffffff8004935b3f
  kernel`_Xmach_ports_register(InHeadP=0xffffffff800b2c297c,
  OutHeadP=0xffffffff800e5c4b90) + 111 at task_server.c:647 [opt]
9     frame #6: 0xffffffff80048df2c3
  kernel`ipc_kobject_server(request=0xffffffff800b2c2900) + 259 at
  ipc_kobject.c:340 [opt]
10    frame #7: 0xffffffff80048c28f8 kernel`ipc_kmsg_send(kmsg=<unavailable>,
  option=<unavailable>, send_timeout=0) + 184 at ipc_kmsg.c:1443 [opt]
11    frame #8: 0xffffffff80048d26a5 kernel`mach_msg_overwrite_trap(args=
  <unavailable>) + 197 at mach_msg.c:474 [opt]
12    frame #9: 0xffffffff80049b8eca
  kernel`mach_call_munger64(state=0xffffffff800f7e6540) + 410 at bsd_i386.c:560
  [opt]
13    frame #10: 0xffffffff80049ecd86 kernel`hndl_mach_scall64 + 22

```

简单的分析一下调用栈

`_Xmach_ports_register` 就是 `taskServer.c` 中对应的函数。

出问题的最关键的是 #4，#5 两个栈。

通过分析 `mach_ports_register` 函数的源码

```

1
2    ...
3
4    for (i = 0; i < TASK_PORT_REGISTER_MAX; i++) {
5        ipc_port_t old;
6
7        old = task->itk_registered[i];
8        task->itk_registered[i] = ports[i];
9        ports[i] = old;
10    }
11
12    itk_unlock(task);
13
14    for (i = 0; i < TASK_PORT_REGISTER_MAX; i++)
15        if (IP_VALID(ports[i]))
16            ipc_port_release_send(ports[i]); <--#5b崩溃的地方
17
18    if (portsCnt != 0)
19        kfree(memory, <--释放memory
20            (vm_size_t) (portsCnt * sizeof(mach_port_t)));
21    ...

```

这里是对 `ports` 的数组中的参数调用 `ipc_port_release_send`，出发的崩溃。

查看 `ports` 中的值，如下，

```

1 (lldb) f 4
2 kernel was compiled with optimization - stepping may behave oddly; variables
  may not be available.
3 frame #4: 0xfffff80048e22d3 kernel`mach_ports_register(task=<unavailable>,
  memory=0xfffff800aad4270, portsCnt=3) + 547 at ipc_tt.c:1097 [opt]
4 (lldb) p ports
5 (ipc_port_t [3]) $0 = {
6     [0] = 0xfffff800b890680
7     [1] = 0xdeadbeefdeadbeef
8     [2] = 0x6c7070612e6d6f63
9 }

```

因为源码中会使用 `kfree` 去释放 `memory`，接下来就去动态的调试吧。

## 2.2 动态调试

### 2.2.1 mach\_ports\_register

因为 `mach_ports_register` 这个函数在一些其他流程中都会有调用，如果直接在内核中的 `mach_ports_register` 下断点，会有很多其他的调用会被断到，这里我的做法是先用 `lldb` 启动 `r3gister` 程序并断在 `mach_ports_register` 处，并运行。

```
1  → lldb r3gister
2  (lldb) target create "r3gister"
3  Current executable set to 'r3gister' (x86_64).
4  (lldb) b mach_ports_register
5  Breakpoint 1: 2 locations.
6  (lldb) r
7  Process 425 launched: '/Users/mrh/mach_port_register/r3gister' (x86_64)
8  Process 425 stopped
9  * thread #1: tid = 0x10fd, 0x00000001000012a7
   r3gister`mach_ports_register(target_task=259,
   init_port_set=0x00007fff5fbffa98, init_port_setCnt=3) + 39 at taskUser.c:690,
   queue = 'com.apple.main-thread', stop reason = breakpoint 1.1
10     frame #0: 0x00000001000012a7 r3gister`mach_ports_register(target_task=259,
   init_port_set=0x00007fff5fbffa98, init_port_setCnt=3) + 39 at taskUser.c:690
11     687         Reply Out;
12     688     } Mess;
13     689
14     -> 690     Request *InP = &Mess.In;
15     691     Reply *Out0P = &Mess.Out;
16     692
17     693     mach_msg_return_t msg_result;
```

当已经断在这里的时候，在内核上下断点。

```
1  (lldb) b mach_ports_register
2  Breakpoint 1: where = kernel.development`mach_ports_register + 40 at
   ipc_tt.c:1060, address = 0xffffffff8009686568
3  (lldb) c
4  Process 1 resuming
```

在内核的断点设置成功后，继续执行 `r3gister`，就会触发内核中的断点。

```

1 (lldb) bt
2 * thread #1: tid = 0x0001, 0xffffffff8009686568
   kernel.development`mach_ports_register(task=0xffffffff8012933640,
   memory=0xffffffff800f7da660, portsCnt=3) + 40 at ipc_tt.c:1060, stop reason =
   breakpoint 1.1
3   * frame #0: 0xffffffff8009686568
   kernel.development`mach_ports_register(task=0xffffffff8012933640,
   memory=0xffffffff800f7da660, portsCnt=3) + 40 at ipc_tt.c:1060 [opt]
4   frame #1: 0xffffffff80096e43ff
   kernel.development`_Xmach_ports_register(InHeadP=0xffffffff8013c7937c,
   OutHeadP=0xffffffff8014efaf90) + 111 at task_server.c:647 [opt]
5   frame #2: 0xffffffff8009683443
   kernel.development`ipc_kobject_server(request=0xffffffff8013c79300) + 259 at
   ipc_kobject.c:340 [opt]
6   frame #3: 0xffffffff800965ef03 kernel.development`ipc_kmsg_send(kmsg=
   <unavailable>, option=<unavailable>, send_timeout=0) + 211 at ipc_kmsg.c:1443
   [opt]
7   frame #4: 0xffffffff8009675985
   kernel.development`mach_msg_overwrite_trap(args=<unavailable>) + 197 at
   mach_msg.c:474 [opt]
8   frame #5: 0xffffffff800977f000
   kernel.development`mach_call_munger64(state=0xffffffff801278eb60) + 480 at
   bsd_i386.c:560 [opt]
9   frame #6: 0xffffffff80097b4de6 kernel.development`hndl_mach_scall64 + 22

```

## 2.2.2 memory的zone分析

通过 `lldb` 调试器查看 `memory` 处的内存，如下。

```

1 (lldb) memory read --format x --size 8 memory
2 0xffffffff800f7da660: 0xffffffff80140e7580 0xdeadbeefdeadbeef
3 0xffffffff800f7da670: 0xffffffff800f7dab00 0xfacadea23d1ec085
4 0xffffffff800f7da680: 0x0000000000000000 0xfffffffff00000000
5 0xffffffff800f7da690: 0x0000000000000000 0x0000000000001000

```

通过一点小技巧可以查看 `memory` 是在哪一个 `zone` 上分配的，也可以继续跟踪代码，在后面的 `kfree` 中调用 `zfree` 函数的流程中会出现相关的转换代码。

```

1 ...
2 if (zone->use_page_list) {
3     struct zone_page_metadata *page_meta = get_zone_page_metadata((struct
   zone_free_element *)addr);
4     if (zone != page_meta->zone) {
5     ...

```

其实就是 `get_zone_page_metadata` 这个函数的实现了，这里的 `addr` 就是 `memory`。

```

1  (lldb) p *(zone_page_metadata*)0xffffffff80140e7000
2  (zone_page_metadata) $1 = {
3      pages = {
4          next = 0xffffffff8012db4000
5          prev = 0xffffffff801109f000
6      }
7      elements = 0xffffffff80140e76d0
8      zone = 0xffffffff800f480ba0
9      alloc_count = 12
10     free_count = 1
11 }

```

在查看 `zone` 的具体数据，可以得知 `memory` 被分配在哪个 `zone` 当中。

```

1  ...
2  zone_name = 0xffffffff8009d35bac "kalloc.16"
3  ...

```

可以得知，因为 `port` 的个数被设置为1个，所以只需要一个指针，而在调用 `kfree` 的时候，`kfree` 的 `size` 是3个指针的长度，所以是试图在 `kalloc.24` 中释放内存,这就会造成错误的 `kfree`，但是苹果有一段神奇的代码，尝试修复这个问题。

```

1  if (zone->use_page_list) {
2      struct zone_page_metadata *page_meta = get_zone_page_metadata((struct
zone_free_element *)addr);
3      if (zone != page_meta->zone) {
4          /*
5           * Something bad has happened. Someone tried to zfree a pointer
but the metadata says it is from
6           * a different zone (or maybe it's from a zone that doesn't use
page free lists at all). We can repair
7           * some cases of this, if:
8           * 1) The specified zone had use_page_list, and the true zone
also has use_page_list set. In that case
9           *     we can swap the zone_t
10          * 2) The specified zone had use_page_list, but the true zone
does not. In this case page_meta is garbage,
11          *     and dereferencing page_meta->zone might panic.
12          * To distinguish the two, we enumerate the zone list to match
it up.
13          * We do not handle the case where an incorrect zone is passed
that does not have use_page_list set,
14          * even if the true zone did have this set.
15          */
16          zone_t fixed_zone = NULL;
17          int fixed_i, max_zones;
18
19          simple_lock(&all_zones_lock);
20          max_zones = num_zones;
21          fixed_zone = first_zone;
22          simple_unlock(&all_zones_lock);
23
24          for (fixed_i=0; fixed_i < max_zones; fixed_i++, fixed_zone =
fixed_zone->next_zone) {
25              if (fixed_zone == page_meta->zone && fixed_zone-
>use_page_list) {
26                  /* we can fix this */
27                  printf("Fixing incorrect zfree from zone %s to zone
%s\n", zone->zone_name, fixed_zone->zone_name);
28                  zone = fixed_zone;
29                  break;
30              }
31          }
32      }
33  }

```

用代码修复数据结构的错误本身就是一件很危险的事情，而这里更危险的是如果不能修复这个错误的话，代码没有任何报错或提示，这里就会有很多隐患。

## 2.2.3 堆内存简单分析

这里简单的分析一下内核中堆的数据结构，写到这里的时候我重启了一次虚拟机和调试器，所以地址和之前会对不上。

这一次看到的 `kalloc.16` 的 `zone` 如下。

```
1  p *(zone*)0xffffffff80213caea0
2  (zone) $5 = {
3      free_elements = 0x0000000000000000
4      pages = {
5          any_free_foreign = {
6              next = 0xffffffff80213caea8
7              prev = 0xffffffff80213caea8
8          }
9          all_free = {
10             next = 0xffffffff80213caeb8
11             prev = 0xffffffff80213caeb8
12         }
13         intermediate = {
14             next = 0xffffffff8025e06000
15             prev = 0xffffffff8025706000
16         }
17         all_used = {
18             next = 0xffffffff8025226000
19             prev = 0xffffffff8025d29000
20         }
21     }
22     count = 29951
23     countfree = 37
24     lock_attr = (lck_attr_val = 0)
25     lock = {
26         lck_mtx_sw = {
27             lck_mtxd = {
28                 lck_mtxd_owner = 0
29                 = {
30                     = {
31                         lck_mtxd_waiters = 0
32                         lck_mtxd_pri = 0
33                         lck_mtxd_ilocked = 0
34                         lck_mtxd_mlocked = 0
35                         lck_mtxd_promoted = 0
36                         lck_mtxd_spin = 0
37                         lck_mtxd_is_ext = 0
38                         lck_mtxd_pad3 = 0
39                     }
40                     lck_mtxd_state = 0
41                 }
42                 lck_mtxd_pad32 = 4294967295
```



```

43     }
44     lck_mtxi = {
45         lck_mtxi_ptr = 0x0000000000000000
46         lck_mtxi_tag = 0
47         lck_mtxi_pad32 = 4294967295
48     }
49 }
50 }
51 lock_ext = {
52     lck_mtx = {
53         lck_mtx_sw = {
54             lck_mtxd = {
55                 lck_mtxd_owner = 0
56                 = {
57                     = {
58                         lck_mtxd_waiters = 0
59                         lck_mtxd_pri = 0
60                         lck_mtxd_ilocked = 0
61                         lck_mtxd_mlocked = 0
62                         lck_mtxd_promoted = 0
63                         lck_mtxd_spin = 0
64                         lck_mtxd_is_ext = 0
65                         lck_mtxd_pad3 = 0
66                     }
67                     lck_mtxd_state = 0
68                 }
69                 lck_mtxd_pad32 = 0
70             }
71             lck_mtxi = {
72                 lck_mtxi_ptr = 0x0000000000000000
73                 lck_mtxi_tag = 0
74                 lck_mtxi_pad32 = 0
75             }
76         }
77     }
78     lck_mtx_grp = 0x0000000000000000
79     lck_mtx_attr = 0
80     lck_mtx_pad1 = 0
81     lck_mtx_deb = (type = 0, pad4 = 0, pc = 0, thread = 0)
82     lck_mtx_stat = 0
83     lck_mtx_pad2 = ([0] = 0, [1] = 0)
84 }
85 cur_size = 479808
86 max_size = 531441
87 elem_size = 16
88 alloc_size = 4096
89 page_count = 119
90 sum_count = 235587
91 exhaustible = 0

```

```
92     collectable = 1
93     expandable = 1
94     allows_foreign = 0
95     doing_alloc_without_vm_priv = 0
96     doing_alloc_with_vm_priv = 0
97     waiting = 0
98     async_pending = 0
99     zleak_on = 0
100    caller_acct = 0
101    doing_gc = 0
102    noencrypt = 0
103    no_callout = 0
104    async_prio_refill = 0
105    gzalloc_exempt = 0
106    alignment_required = 0
107    use_page_list = 1
108    _reserved = 0
109    index = 12
110    next_zone = 0xffffffff80213ca120
111    zone_name = 0xffffffff801ef3af0d "kalloc.16"
112    zleak_capture = 0
113    zp_count = 0
114    prio_refill_watermark = 0
115    zone_replenish_thread = 0x0000000000000000
116    gz = {
117        gzfc_index = 0
118        gzfc = 0xdeadbeefdeadbeef
119    }
120 }
```

这里就主要的分析一下 `page` 和 `page` 内的内存的分布。

```

1 p *(zone*)0xffffffff80213caea0
2 (zone) $5 = {
3     free_elements = 0x0000000000000000
4     pages = {
5         any_free_foreign = {
6             next = 0xffffffff80213caea8
7             prev = 0xffffffff80213caea8
8         }
9         all_free = {
10             next = 0xffffffff80213caeb8
11             prev = 0xffffffff80213caeb8
12         }
13         intermediate = {
14             next = 0xffffffff8025e06000
15             prev = 0xffffffff8025706000
16         }
17         all_used = {
18             next = 0xffffffff8025226000
19             prev = 0xffffffff8025d29000
20         }
21     }
22     ...

```

简单的看一下4个 `pages` 的队列中的 `intermediate`，在这个队列中的 `page` 里都会有一些未被使用的内存,通过 `pages` 里面的 `next` 和 `prev` 构成了一个双向链表，如下所示。

```

1  (lldb) p *(zone_page_metadata*)0xffffffff8025e06000
2  (zone_page_metadata) $6 = {
3      pages = {
4          next = 0xffffffff8025548000
5          prev = 0xffffffff80213caec8
6      }
7      elements = 0xffffffff8025e06750
8      zone = 0xffffffff80213caea0
9      alloc_count = 252
10     free_count = 1
11 }
12
13 (lldb) p *(zone_page_metadata*)0xffffffff8025548000
14 (zone_page_metadata) $12 = {
15     pages = {
16         next = 0xffffffff8025c96000
17         prev = 0xffffffff8025e06000
18     }
19     elements = 0xffffffff8025548d50
20     zone = 0xffffffff80213caea0
21     alloc_count = 252
22     free_count = 1
23 }
24
25 (lldb) p *(zone_page_metadata*)0xffffffff8025c96000
26 (zone_page_metadata) $13 = {
27     pages = {
28         next = 0xffffffff8025cb3000
29         prev = 0xffffffff8025548000
30     }
31     elements = 0xffffffff8025c968e0
32     zone = 0xffffffff80213caea0
33     alloc_count = 252
34     free_count = 5
35 }

```

`elements` 就是第一个可以 `alloc` 的内存，通过lldb观察内存布局

```

1 (lldb) memory read --format x --size 8 0xffffffff8025c968e0
2 0xffffffff8025c968e0: 0xffffffff8025c96670 0xfacade04d7b687dd
3
4 (lldb) memory read --format x --size 8 0xffffffff8025c96670
5 0xffffffff8025c96670: 0xffffffff8025c96680 0xfacade04d7b6872d
6
7 (lldb) memory read --format x --size 8 0xffffffff8025c96680
8 0xffffffff8025c96680: 0xffffffff8025c96a40 0xfacade04d7b68bed
9
10 (lldb) memory read --format x --size 8 0xffffffff8025c96680
11 0xffffffff8025c96680: 0xffffffff8025c96a40 0xfacade04d7b68bed

```

`freeelement` 是通过单向链表随机的串联在 `page` 中，在前面[iOS 10 Kernel Heap Revisited](#)中提到的。

可以看到前面的8个字节控制链表的，后面8个字节是 `freeelement` 的存储空间，`0xfacade` 就是堆中的cookies。

```

1 zp_poisoned_cookie &= 0x000000FFFFFFFFFF;
2 zp_poisoned_cookie |= 0x0535210000000000; /* 0xFACADE */

```

了解了 `freeelement` 的内存布局，再看一看已经被分配了的内存，也就是 `memory`。

```

1 (lldb) memory read --format x --size 8 memory
2 0xffffffff8025e06300: 0xffffffff8029828430 0xdeadbeefdeadbeef

```

阅读源码中的 `zalloc_internal` 函数的实现，可以得知，在 `kalloc.16` 的堆中分配申请内存时，会将申请出来的内存会被写入 `0xdeadbeefdeadbeef`。

```

1 vm_offset_t *primary = (vm_offset_t *) addr; //addr == memory
2 vm_offset_t *backup  = get_backup_ptr(inner_size, primary);
3
4 *primary = ZP_POISON;
5 *backup  = ZP_POISON;

```

## 2.2.4 ipc\_port\_release\_send

在导致崩溃的函数处下断点

```

1  (lldb) b ipc_tt.c :1097
2  Breakpoint 2: where = kernel.development`mach_ports_register + 521 at
   ipc_tt.c:1097, address = 0xffffffff801e886749
3  (lldb) c
4  Process 1 resuming
5
6  Process 1 stopped
7  * thread #2: tid = 0x0002, 0xffffffff801e886749
   kernel.development`mach_ports_register(task=<unavailable>,
   memory=0xffffffff8025e06300, portsCnt=3) + 521 at ipc_tt.c:1097, stop reason =
   breakpoint 2.1
8      frame #0: 0xffffffff801e886749 kernel.development`mach_ports_register(task=
   <unavailable>, memory=0xffffffff8025e06300, portsCnt=3) + 521 at ipc_tt.c:1097
   [opt]
9
10 (lldb) p ports
11 (ipc_port_t [3]) $15 = {
12     [0] = 0xffffffff80279d8190
13     [1] = 0x0000000000000000
14     [2] = 0x0000000000000000
15 }

```

发现 `ports` 的值只有 `ports[0]` 有值，这是因为这里的 `ports` 是从旧的 `port` 中替换来的

```

1      /*
2       * Replace the old send rights with the new.
3       * Release the old rights after unlocking.
4       */
5
6      for (i = 0; i < TASK_PORT_REGISTER_MAX; i++) {
7          ipc_port_t old;
8
9          old = task->itk_registered[i];
10         task->itk_registered[i] = ports[i];
11         ports[i] = old;
12     }

```

据悉执行后，`r3gister` 会再次被断住，第二次调用 `mach_ports_register` 后，内核断点，再看 `ports`，如下

```

1  * thread #2: tid = 0x0002, 0xffffffff801e886749
   kernel.development`mach_ports_register(task=<unavailable>,
   memory=0xffffffff80253d3e70, portsCnt=3) + 521 at ipc_tt.c:1097, stop reason =
   breakpoint 2.1
2      frame #0: 0xffffffff801e886749 kernel.development`mach_ports_register(task=
   <unavailable>, memory=0xffffffff80253d3e70, portsCnt=3) + 521 at ipc_tt.c:1097
   [opt]
3  (lldb) p ports
4  (ipc_port_t [3]) $16 = {
5      [0] = 0xffffffff8029828430
6      [1] = 0xdeadbeefdeadbeef
7      [2] = 0xffffffff80252f1460
8  }

```

从而导致在服务器的后续代码中触发了崩溃

```

1      for (i = 0; i < TASK_PORT_REGISTER_MAX; i++)
2          if (IP_VALID(ports[i]))
3              ipc_port_release_send(ports[i]); //<--第二个ports就是
   0xdeadbeefdeadbeef
4
5      /*
6       * Now that the operation is known to be successful,
7       * we can free the memory.
8       */
9
10     if (portsCnt != 0)
11         kfree(memory,
12             (vm_size_t) (portsCnt * sizeof(mach_port_t)));

```

## 0x03 小结

这里只分析了 POC 的触发，如果修改 `mach_ports_register` 的参数，将 `port` 的个数改为2个，`port[1]`就会变成 `0x0000000000000000`，从而避免了对 `0xdeadbeefdeadbeef` 调用函数出发崩溃，而且zone会变成iports的一个专用的 `zone`，而不是 `kalloc.16`，所以这个漏洞值得研究的地方还有很多。希望本文能为大家继续研究这个漏洞提供一些帮助；-）。

## 参考

- 1、[OS X/iOS multiple memory safety issues in mach\\_ports\\_register](#)
- 2、[iOS 10 Kernel Heap Revisited](#)

